

Experiment 3: Classification Algorithm Comparisons

Name: Tarun Suresh

Register Number: 3122237001055

Date: 12/08/25

Department of Computer Science and Engineering
SSN College of Engineering

Academic Year: 2024–2025

```
#mounting drive
from google.colab import drive
drive.mount('/content/drive')
```

↻ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount(

```
#importing libraries for classification
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
#importing dataset
df = pd.read_csv('/content/drive/MyDrive/spambase_csv.csv')
df
```

↻

	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	word_freq...
0	0.00	0.64	0.64	0.0	
1	0.21	0.28	0.50	0.0	
2	0.06	0.00	0.71	0.0	
3	0.00	0.00	0.00	0.0	
4	0.00	0.00	0.00	0.0	
...	
4596	0.31	0.00	0.62	0.0	
4597	0.00	0.00	0.00	0.0	
4598	0.30	0.00	0.30	0.0	
4599	0.96	0.00	0.00	0.0	
4600	0.00	0.00	0.65	0.0	

4601 rows × 58 columns

```
#performing eda
missing_values = df.isna().sum()
```

```
print(missing_values)
```

```
#dealing with missing values
# Fill numerical columns with mean
df.fillna(df.mean(), inplace=True)
```

```
⇒ word_freq_make 0
word_freq_address 0
word_freq_all 0
word_freq_3d 0
word_freq_our 0
word_freq_over 0
word_freq_remove 0
word_freq_internet 0
word_freq_order 0
word_freq_mail 0
word_freq_receive 0
word_freq_will 0
word_freq_people 0
word_freq_report 0
word_freq_addresses 0
word_freq_free 0
word_freq_business 0
word_freq_email 0
word_freq_you 0
word_freq_credit 0
word_freq_your 0
word_freq_font 0
word_freq_000 0
word_freq_money 0
word_freq_hp 0
word_freq_hpl 0
word_freq_george 0
word_freq_650 0
word_freq_lab 0
word_freq_labs 0
word_freq_telnet 0
word_freq_857 0
word_freq_data 0
word_freq_415 0
word_freq_85 0
word_freq_technology 0
word_freq_1999 0
word_freq_parts 0
word_freq_pm 0
word_freq_direct 0
word_freq_cs 0
word_freq_meeting 0
word_freq_original 0
word_freq_project 0
word_freq_re 0
word_freq_edu 0
```

```

word_freq_table      0
word_freq_conference 0
char_freq_%3B        0
char_freq_%28        0
char_freq_%5B        0
char_freq_%21        0
char_freq_%24        0
char_freq_%23        0
capital_run_length_average 0
capital_run_length_longest 0
capital_run_length_total 0
class                0
dtype: int64

```

```
# Check for outliers visually using boxplots
```

```

plt.figure(figsize=(20, 15))
sns.boxplot(data=df)
plt.title('Boxplot of Scaled Features')
plt.xticks(rotation=90)
plt.show()

```

```
# Check for outliers programmatically using IQR
```

```

Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1

```

```
# Define outlier bounds
```

```

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

```

```
# Find outliers
```


```


outliers = ((df < lower_bound) | (df > upper_bound)).sum()
print("\nNumber of outliers per column (IQR method):")
print(outliers[outliers > 0])

```

```
#removing outliers
```

```
#df = df[~((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).any(axis=1)]
```





Number of outliers per column (IQR method):

word_freq_make	1053
word_freq_address	898
word_freq_all	338
word_freq_3d	47
word_freq_our	501
word_freq_over	999
word_freq_remove	807
word_freq_internet	824
word_freq_order	773
word_freq_mail	852
word_freq_receive	709
word_freq_will	270
word_freq_people	852
word_freq_report	357
word_freq_addresses	336
word_freq_free	957
word_freq_business	963
word_freq_email	1038
word_freq_you	75
word_freq_credit	424
word_freq_your	229
word_freq_font	117
word_freq_000	679
word_freq_money	735
word_freq_hp	1090
word_freq_hpl	811
word_freq_george	780
word_freq_650	463
word_freq_lab	372
word_freq_labs	469
word_freq_telnet	293
word_freq_857	205
word_freq_data	405


word_freq_data	100
word_freq_415	215
word_freq_85	485
word_freq_technology	599
word_freq_1999	829
word_freq_parts	83
word_freq_pm	384
word_freq_direct	453
word_freq_cs	148
word_freq_meeting	341
word_freq_original	375
word_freq_project	327
word_freq_re	1001
word_freq_edu	517
word_freq_table	63
word_freq_conference	203
char_freq_%3B	790
char_freq_%28	296
char_freq_%5B	529
char_freq_%21	411
char_freq_%24	811
char_freq_%23	750
capital_run_length_average	363
capital_run_length_longest	463
capital_run_length_total	550
dtype:	int64

```
y = df['class']  
df = df.drop('class', axis=1)  
  
#using standard scaler on data for gaussianNB  
from sklearn.preprocessing import StandardScaler  
  
scaler = StandardScaler()  
df_scaled = scaler.fit_transform(df)
```

```
#normalise the values for gaussianNB
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
df_scaled = scaler.fit_transform(df)
```

df




	word_freq_make	word_freq_address	word_freq_all	word_freq_3d	word_freq...
0	0.00	0.64	0.64	0.0	
1	0.21	0.28	0.50	0.0	
2	0.06	0.00	0.71	0.0	
3	0.00	0.00	0.00	0.0	
4	0.00	0.00	0.00	0.0	
...	
4596	0.31	0.00	0.62	0.0	
4597	0.00	0.00	0.00	0.0	
4598	0.30	0.00	0.30	0.0	
4599	0.96	0.00	0.00	0.0	
4600	0.00	0.00	0.65	0.0	

4601 rows x 57 columns

```
#binarize dataset from original data set for NaiveBayes Bernoulli
from sklearn.preprocessing import Binarizer
```

```
binarizer = Binarizer(threshold=0.5)
df_binarized = binarizer.fit_transform(df)
```

df_binarized



```
array([[0., 1., 1., ..., 1., 1., 1.],
       [0., 0., 0., ..., 1., 1., 1.],
       [0., 0., 1., ..., 1., 1., 1.],
       ...,
       [0., 0., 0., ..., 1., 1., 1.],
       [1., 0., 0., ..., 1., 1., 1.],
       [0., 0., 1., ..., 1., 1., 1.]])
```


We will be using a normalised dataset for Naive Bayes Gaussian distribution as it expects features to be normalised during fit and predict.

```
#splitting normalised dataset and performing NaiveBayes Gaussian
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

Xg_train, Xg_test, yg_train, yg_test = train_test_split(df_scaled, y, test_size=0.2)

gNBmodel = GaussianNB()
gNBmodel.fit(Xg_train, yg_train)
yg_pred = gNBmodel.predict(Xg_test)

#evaluating Accuracy, Precision, Recall, F1-score for GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

accuracy_g = accuracy_score(yg_test, yg_pred)
precision_g = precision_score(yg_test, yg_pred)
recall_g = recall_score(yg_test, yg_pred)
f1_g = f1_score(yg_test, yg_pred)

print("Accuracy:", accuracy_g)
print("Precision:", precision_g)
print("Recall:", recall_g)
print("F1-score:", f1_g)



🔗 Accuracy: 0.8219326818675353
Precision: 0.7233201581027668
Recall: 0.9384615384615385
F1-score: 0.8169642857142857


#displaying confusion matrix
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
cm = confusion_matrix(yg_test, yg_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')
plt.show()

#displaying roc curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(yg_test, yg_pred)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```





As for NB Multinomial we will be using the original dataset as it depends on raw counts/frequency of features.

```
#k-fold for NB gaussian
from sklearn.model_selection import cross_val_score
from sklearn.naive_bayes import GaussianNB

gNBmodel = GaussianNB()
scores = cross_val_score(gNBmodel, df_scaled, y, cv=5)
print("Cross-validation scores:", scores)
print("Mean accuracy:", scores.mean())
```

```
↔ Cross-validation scores: [0.85124864 0.86630435 0.85434783 0.84347826 0.695
Mean accuracy: 0.8222062502950479
```

```
#splitting original dataset and performing NaiveBayes Multinomial
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score

Xm_train, Xm_test, ym_train, ym_test = train_test_split(df, y, test_size=0.2, r

mNBmodel = MultinomialNB()
mNBmodel.fit(Xm_train, ym_train)
ym_pred = mNBmodel.predict(Xm_test)

accuracy_m = accuracy_score(ym_test, ym_pred)
precision_m = precision_score(ym_test, ym_pred)
recall_m = recall_score(ym_test, ym_pred)
f1_m = f1_score(ym_test, ym_pred)

print("Accuracy:", accuracy_g)
print("Precision:", precision_g)
print("Recall:", recall_g)
print("F1-score:", f1_g)
```



```
➦ Accuracy: 0.8219326818675353
Precision: 0.7233201581027668
Recall: 0.9384615384615385
F1-score: 0.8169642857142857
```


```
#displaying confusion matrix
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
cm = confusion_matrix(ym_test, ym_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

```
#displaying roc curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(ym_test, ym_pred)
roc_auc = auc(fpr, tpr)
```

```
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
```

```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```





```
#k-fold for NB multinomial
from sklearn.model_selection import cross_val_score
from sklearn.naive_bayes import MultinomialNB
```

```
mNBmodel = MultinomialNB()
scores = cross_val_score(mNBmodel, df, y, cv=5)
print("Cross-validation scores:", scores)
print("Mean accuracy:", scores.mean())
```

```
↔ Cross-validation scores: [0.79261672 0.81847826 0.81521739 0.78586957 0.696
Mean accuracy: 0.7817842137563139
```

For NB bernoulli we will use Binariser to convert the data to binary

```
#train test split and NB Bernoulli with binarised data
from sklearn.naive_bayes import BernoulliNB
from sklearn.metrics import accuracy_score

Xb_train, Xb_test, yb_train, yb_test = train_test_split(df_binarized, y, test_s

bNBmodel = BernoulliNB()
bNBmodel.fit(Xb_train, yb_train)
yb_pred = bNBmodel.predict(Xb_test)

accuracy_b = accuracy_score(yb_test, yb_pred)
precision_b = precision_score(yb_test, yb_pred)
recall_b = recall_score(yb_test, yb_pred)
f1_b = f1_score(yb_test, yb_pred)

print("Accuracy:", accuracy_b)
print("Precision:", precision_b)
print("Recall:", recall_b)
print("F1-score:", f1_b)
```



```
➦ Accuracy: 0.8957654723127035
Precision: 0.8712121212121212
Recall: 0.8846153846153846
F1-score: 0.8778625954198473
```


```
#displaying confusion matrix
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
cm = confusion_matrix(yb_test, yb_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

```
#displaying roc curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(yb_test, yb_pred)
roc_auc = auc(fpr, tpr)
```

```
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
```

```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```





```
#applying K-fold Cross validation for NB bernoulli
from sklearn.model_selection import cross_val_score
from sklearn.naive_bayes import BernoulliNB
```

```
bNBmodel = BernoulliNB()
scores = cross_val_score(bNBmodel, df_binarized, y, cv=5)
print("Cross-validation scores:", scores)
print("Mean accuracy:", scores.mean())
```

```
⇒ Cross-validation scores: [0.90662324 0.91086957 0.92391304 0.93695652 0.744
Mean accuracy: 0.8845855166879101
```

```
#mounting drive
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
#importing libraries for classification
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#importing dataset
df = pd.read_csv('/content/drive/MyDrive/spambase_csv.csv')
df

#performing eda
missing_values = df.isna().sum()
print(missing_values)

#dealing with missing values
# Fill numerical columns with mean
df.fillna(df.mean(), inplace=True)

# Check for outliers visually using boxplots
plt.figure(figsize=(20, 15))
sns.boxplot(data=df)
plt.title('Boxplot of Scaled Features')
plt.xticks(rotation=90)
plt.show()

# Check for outliers programmatically using IQR
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1

# Define outlier bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Find outliers
outliers = ((df < lower_bound) | (df > upper_bound)).sum()
```

```
print("\nNumber of outliers per column (IQR method):")
print(outliers[outliers > 0])
```

```
#removing outliers
```


```
#df = df[~((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).any(axis=1)]
```

```
word_freq_make 0
word_freq_address 0
word_freq_all 0
word_freq_3d 0
word_freq_our 0
word_freq_over 0
word_freq_remove 0
word_freq_internet 0
word_freq_order 0
word_freq_mail 0
word_freq_receive 0
word_freq_will 0
word_freq_people 0
word_freq_report 0
word_freq_addresses 0
word_freq_free 0
word_freq_business 0
word_freq_email 0
word_freq_you 0
word_freq_credit 0
word_freq_your 0
word_freq_font 0
word_freq_000 0
word_freq_money 0
word_freq_hp 0
word_freq_hpl 0
word_freq_george 0
word_freq_650 0
word_freq_lab 0
word_freq_labs 0
word_freq_telnet 0
word_freq_857 0
word_freq_data 0
word_freq_415 0
word_freq_85 0
word_freq_technology 0
word_freq_1999 0
word_freq_parts 0
word_freq_pm 0
word_freq_direct 0
word_freq_cs 0
word_freq_meeting 0
word_freq_original 0
word_freq_project 0
word_freq_re 0
word_freq_edu 0
```

```

word_freq_table          0
word_freq_conference     0
char_freq_%3B            0
char_freq_%28            0
char_freq_%5B            0
char_freq_%21            0
char_freq_%24            0
char_freq_%23            0
capital_run_length_average 0
capital_run_length_longest 0
capital_run_length_total  0
class                    0
dtype: int64

```



Number of outliers per column (IQR method):

```

word_freq_make          1053
word_freq_address       898
word_freq_all           338
word_freq_3d            47
word_freq_our           501
word_freq_over          999
word_freq_remove        807

```

word_freq_internet	824
word_freq_order	773
word_freq_mail	852
word_freq_receive	709
word_freq_will	270
word_freq_people	852
word_freq_report	357
word_freq_addresses	336
word_freq_free	957
word_freq_business	963
word_freq_email	1038
word_freq_you	75
word_freq_credit	424
word_freq_your	229
word_freq_font	117
word_freq_000	679
word_freq_money	735
word_freq_hp	1090
word_freq_hpl	811
word_freq_george	780
word_freq_650	463
word_freq_lab	372
word_freq_labs	469
word_freq_telnet	293
word_freq_857	205
word_freq_data	405
word_freq_415	215
word_freq_85	485
word_freq_technology	599
word_freq_1999	829
word_freq_parts	83
word_freq_pm	384
word_freq_direct	453
word_freq_cs	148
word_freq_meeting	341
word_freq_original	375
word_freq_project	327
word_freq_re	1001
word_freq_edu	517
word_freq_table	63
word_freq_conference	203
char_freq_%3B	790
char_freq_%28	296
char_freq_%5B	529
char_freq_%21	411
char_freq_%24	811
char_freq_%23	750
capital_run_length_average	363
capital_run_length_longest	463
capital_run_length_total	550
dtype:	int64

```
#dropping target variable
y = df['class']
df = df.drop('class', axis=1)

#using standard scaler on data
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)
```

```
#performing KNN classification on data
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_ma

X_train, X_test, y_train, y_test = train_test_split(df_scaled, y, test_size=0.2

knn_classifier5 = KNeighborsClassifier(n_neighbors=5)
knn_classifier5.fit(X_train, y_train)

y_pred5 = knn_classifier5.predict(X_test)
```

```
#Prediction analysis
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

accuracy = accuracy_score(y_test, y_pred5)
precision = precision_score(y_test, y_pred5)
recall = recall_score(y_test, y_pred5)
f1 = f1_score(y_test, y_pred5)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
```

```
➡ Accuracy: 0.8957654723127035
Precision: 0.9016393442622951
Recall: 0.8461538461538461
F1-score: 0.873015873015873
```

```
#displaying confusion matrix
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
cm = confusion_matrix(y_test, y_pred5)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
#displaying roc curve
```


```

from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred5)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

```





```
#K valye = 1
knn_classifier1 = KNeighborsClassifier(n_neighbors=1)
knn_classifier1.fit(X_train, y_train)

y_pred1 = knn_classifier1.predict(X_test)

#Prediction analysis
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

accuracy = accuracy_score(y_test, y_pred1)
precision = precision_score(y_test, y_pred1)
recall = recall_score(y_test, y_pred1)
f1 = f1_score(y_test, y_pred1)



print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)


#displaying confusion matrix
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
cm = confusion_matrix(y_test, y_pred1)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

```
#displaying roc curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred1)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

```
↔ Accuracy: 0.8957654723127035
Precision: 0.8769230769230769
Recall: 0.8769230769230769
F1-score: 0.8769230769230769
```





```
#K value = 3
knn_classifier = KNeighborsClassifier(n_neighbors=3)
knn_classifier.fit(X_train, y_train)

y_pred = knn_classifier.predict(X_test)

#Prediction analysis
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)

#displaying confusion matrix
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

```

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

```

```


#displaying roc curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)


```

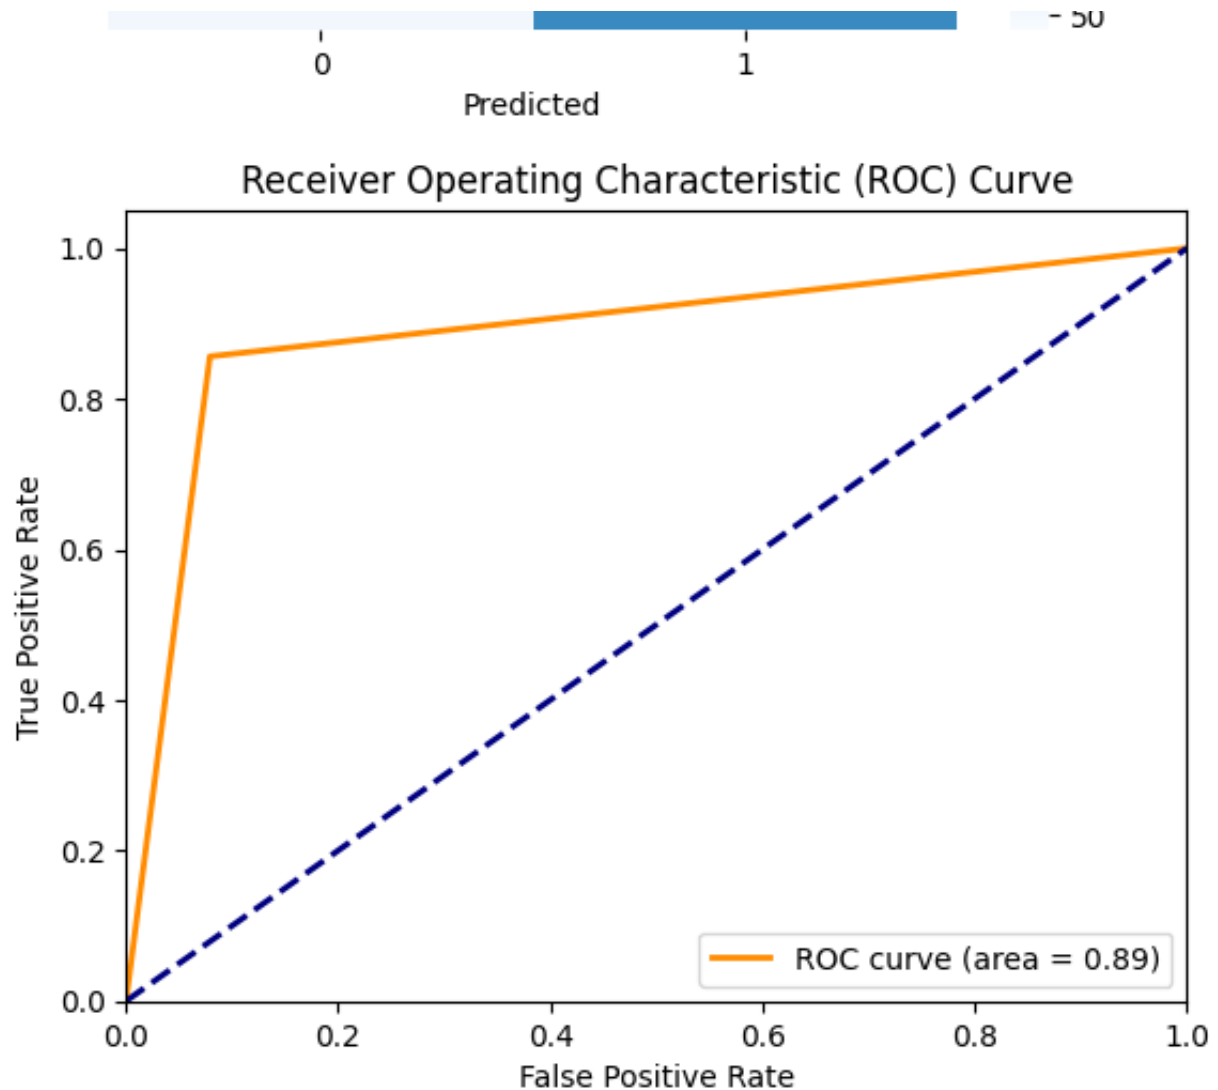
```

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

```

 Accuracy: 0.8935939196525515
Precision: 0.8882978723404256
Recall: 0.8564102564102564
F1-score: 0.8720626631853786





```
#k value = 7
knn_classifier = KNeighborsClassifier(n_neighbors=7)
knn_classifier.fit(X_train, y_train)

y_pred = knn_classifier.predict(X_test)

#Prediction analysis
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("Accuracy:", accuracy)
print("Precision:", precision)
```

```

print("Recall:", recall)
print("F1-score:", f1)

#displaying confusion matrix
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

#displaying roc curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)


plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

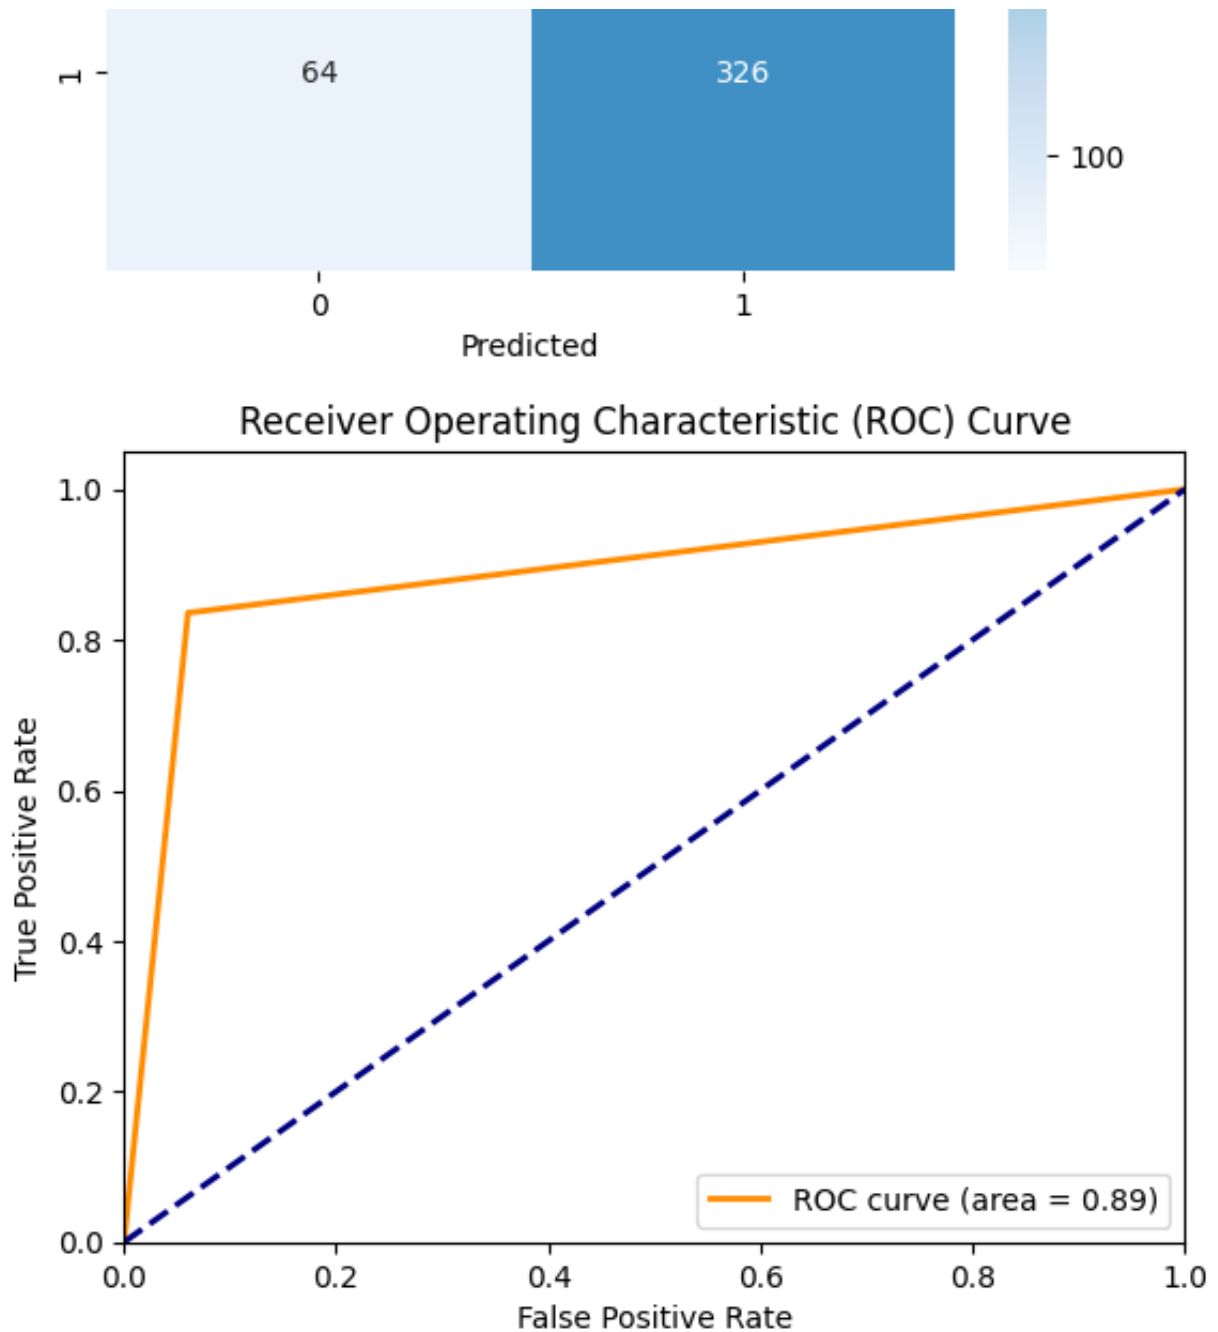
```

```

➡ Accuracy: 0.8957654723127035
Precision: 0.9106145251396648
Recall: 0.8358974358974359
F1-score: 0.8716577540106952

```





```
#Kfold cross validation for 1 neighbour KNN
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier

knn_classifier = KNeighborsClassifier(n_neighbors=1)
cv_scores = cross_val_score(knn_classifier, X_train, y_train, cv=5, scoring='acc')
```

```
#Kfold cross validation for 3 neighbour KNN
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier

knn_classifier = KNeighborsClassifier(n_neighbors=3)
cv_scores = cross_val_score(knn_classifier, X_train, y_train, cv=5, scoring='acc
```

```
#Kfold cross validation for 5 neighbour KNN
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier

knn_classifier = KNeighborsClassifier(n_neighbors=5)
cv_scores = cross_val_score(knn_classifier, X_train, y_train, cv=5, scoring='acc
```

```
#Kfold cross validation for 7 neighbour KNN
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier

knn_classifier = KNeighborsClassifier(n_neighbors=7)
cv_scores = cross_val_score(knn_classifier, X_train, y_train, cv=5, scoring='ac
```



```
#Applying KDtree on the split dataset and performing metrics
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

knn_classifier = KNeighborsClassifier(n_neighbors=5, algorithm='kd_tree')
knn_classifier.fit(X_train, y_train)

y_predt = knn_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_predt)
precision = precision_score(y_test, y_predt)
recall = recall_score(y_test, y_predt)
f1 = f1_score(y_test, y_predt)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
```

➡ Accuracy: 0.8957654723127035
Precision: 0.9016393442622951
Recall: 0.8461538461538461
F1-score: 0.873015873015873

```
#Applying KDball on the split dataset and performing metrics
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

knn_classifier = KNeighborsClassifier(n_neighbors=5, algorithm='ball_tree')
knn_classifier.fit(X_train, y_train)


y_predb = knn_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_predb)
precision = precision_score(y_test, y_predb)
recall = recall_score(y_test, y_predb)
f1 = f1_score(y_test, y_predb)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1-score:", f1)
```

↔ Accuracy: 0.8957654723127035
Precision: 0.9016393442622951
Recall: 0.8461538461538461
F1-score: 0.873015873015873

```
#mounting drive
from google.colab import drive
drive.mount('/content/drive')
```

 Mounted at /content/drive

```
#importing libraries for classification
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#importing dataset
df = pd.read_csv('/content/drive/MyDrive/spambase_csv.csv')
df

#performing eda
missing_values = df.isna().sum()
print(missing_values)

#dealing with missing values
# Fill numerical columns with mean
df.fillna(df.mean(), inplace=True)

# Check for outliers visually using boxplots
plt.figure(figsize=(20, 15))
sns.boxplot(data=df)
plt.title('Boxplot of Scaled Features')
plt.xticks(rotation=90)
plt.show()

# Check for outliers programmatically using IQR
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1

# Define outlier bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Find outliers
outliers = ((df < lower_bound) | (df > upper_bound)).sum()
```

```
print("\nNumber of outliers per column (IQR method):")
print(outliers[outliers > 0])
```

```
#removing outliers
```


```
df = df[~((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).any(axis=1)]
```

```
word_freq_make 0
word_freq_address 0
word_freq_all 0
word_freq_3d 0
word_freq_our 0
word_freq_over 0
word_freq_remove 0
word_freq_internet 0
word_freq_order 0
word_freq_mail 0
word_freq_receive 0
word_freq_will 0
word_freq_people 0
word_freq_report 0
word_freq_addresses 0
word_freq_free 0
word_freq_business 0
word_freq_email 0
word_freq_you 0
word_freq_credit 0
word_freq_your 0
word_freq_font 0
word_freq_000 0
word_freq_money 0
word_freq_hp 0
word_freq_hpl 0
word_freq_george 0
word_freq_650 0
word_freq_lab 0
word_freq_labs 0
word_freq_telnet 0
word_freq_857 0
word_freq_data 0
word_freq_415 0
word_freq_85 0
word_freq_technology 0
word_freq_1999 0
word_freq_parts 0
word_freq_pm 0
word_freq_direct 0
word_freq_cs 0
word_freq_meeting 0
word_freq_original 0
word_freq_project 0
word_freq_re 0
```

```

word_freq_edu          0
word_freq_table        0
word_freq_conference   0
char_freq_%3B          0
char_freq_%28          0
char_freq_%5B          0
char_freq_%21          0
char_freq_%24          0
char_freq_%23          0
capital_run_length_average 0
capital_run_length_longest 0
capital_run_length_total 0
class                  0
dtype: int64

```



Number of outliers per column (IQR method):

```

word_freq_make          1053
word_freq_address       898
word_freq_all           338
word_freq_3d            47
word_freq_our           501
word_freq_over          999

```

word_freq_remove	807
word_freq_internet	824
word_freq_order	773
word_freq_mail	852
word_freq_receive	709
word_freq_will	270
word_freq_people	852
word_freq_report	357
word_freq_addresses	336
word_freq_free	957
word_freq_business	963
word_freq_email	1038
word_freq_you	75
word_freq_credit	424
word_freq_your	229
word_freq_font	117
word_freq_000	679
word_freq_money	735
word_freq_hp	1090
word_freq_hpl	811
word_freq_george	780
word_freq_650	463
word_freq_lab	372
word_freq_labs	469
word_freq_telnet	293
word_freq_857	205
word_freq_data	405
word_freq_415	215
word_freq_85	485
word_freq_technology	599
word_freq_1999	829
word_freq_parts	83
word_freq_pm	384
word_freq_direct	453
word_freq_cs	148
word_freq_meeting	341
word_freq_original	375
word_freq_project	327
word_freq_re	1001
word_freq_edu	517
word_freq_table	63
word_freq_conference	203
char_freq_%3B	790
char_freq_%28	296
char_freq_%5B	529
char_freq_%21	411
char_freq_%24	811
char_freq_%23	750
capital_run_length_average	363
capital_run_length_longest	463
capital_run_length_total	550
dtype:	int64

```
#dropping target variable
y = df['class']
df = df.drop('class', axis=1)

#using standard scaler on data
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)

#performing SVC for each type of kernel [linear,polynomial,RBF,Sigmoid]
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
```

```

from sklearn.metrics import accuracy_score, classification_report, confusion_ma

X_train, X_test, y_train, y_test = train_test_split(df_scaled, y, test_size=0.2)

#performing grid search to find the best parameters for each kernel
from sklearn.model_selection import GridSearchCV

param_grid = {'C': [0.1, 1, 10, 100], 'gamma': ['scale','auto',1, 0.1, 0.01, 0.

svm = SVC()
grid_search = GridSearchCV(svm, param_grid, refit=True, verbose=3)
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
best_score = grid_search.best_score_
print(f"Best Hyperparameters: {best_params}")
print(f"Best Cross-Validation Accuracy: {best_score:.4f}")

```

```

[CV 5/5] END .....C=100, gamma=1, kernel=sigmoid;; score=0.929 total time=
[CV 1/5] END ...C=100, gamma=0.1, kernel=linear;; score=0.897 total time=
[CV 2/5] END ...C=100, gamma=0.1, kernel=linear;; score=0.857 total time=
[CV 3/5] END ...C=100, gamma=0.1, kernel=linear;; score=0.857 total time=
[CV 4/5] END ...C=100, gamma=0.1, kernel=linear;; score=0.929 total time=
[CV 5/5] END ...C=100, gamma=0.1, kernel=linear;; score=0.893 total time=
[CV 1/5] END .....C=100, gamma=0.1, kernel=poly;; score=0.793 total time=
[CV 2/5] END .....C=100, gamma=0.1, kernel=poly;; score=0.786 total time=
[CV 3/5] END .....C=100, gamma=0.1, kernel=poly;; score=0.893 total time=
[CV 4/5] END .....C=100, gamma=0.1, kernel=poly;; score=0.857 total time=
[CV 5/5] END .....C=100, gamma=0.1, kernel=poly;; score=0.893 total time=
[CV 1/5] END .....C=100, gamma=0.1, kernel=rbf;; score=0.897 total time=
[CV 2/5] END .....C=100, gamma=0.1, kernel=rbf;; score=0.821 total time=
[CV 3/5] END .....C=100, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=100, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=100, gamma=0.1, kernel=rbf;; score=0.929 total time=
[CV 1/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.862 total time=
[CV 2/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.893 total time=
[CV 4/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.821 total time=
[CV 5/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.857 total time=
[CV 1/5] END ..C=100, gamma=0.01, kernel=linear;; score=0.897 total time=
[CV 2/5] END ..C=100, gamma=0.01, kernel=linear;; score=0.857 total time=
[CV 3/5] END ..C=100, gamma=0.01, kernel=linear;; score=0.857 total time=
[CV 4/5] END ..C=100, gamma=0.01, kernel=linear;; score=0.929 total time=
[CV 5/5] END ..C=100, gamma=0.01, kernel=linear;; score=0.893 total time=
[CV 1/5] END ....C=100, gamma=0.01, kernel=poly;; score=0.862 total time=
[CV 2/5] END ....C=100, gamma=0.01, kernel=poly;; score=0.929 total time=
[CV 3/5] END ....C=100, gamma=0.01, kernel=poly;; score=0.857 total time=
[CV 4/5] END ....C=100, gamma=0.01, kernel=poly;; score=0.893 total time=
[CV 5/5] END ....C=100, gamma=0.01, kernel=poly;; score=0.893 total time=
[CV 1/5] END .....C=100, gamma=0.01, kernel=rbf;; score=0.897 total time=
[CV 2/5] END .....C=100, gamma=0.01, kernel=rbf;; score=0.893 total time=

```



```
[CV 3/5] END .....C=100, gamma=0.01, kernel=rbf;; score=0.857 total time=
[CV 4/5] END .....C=100, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=100, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.828 total time=
[CV 2/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.821 total time=
[CV 4/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.893 total time=
[CV 5/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.893 total time=
[CV 1/5] END .C=100, gamma=0.001, kernel=linear;; score=0.897 total time=
[CV 2/5] END .C=100, gamma=0.001, kernel=linear;; score=0.857 total time=
[CV 3/5] END .C=100, gamma=0.001, kernel=linear;; score=0.857 total time=
[CV 4/5] END .C=100, gamma=0.001, kernel=linear;; score=0.929 total time=
[CV 5/5] END .C=100, gamma=0.001, kernel=linear;; score=0.893 total time=
[CV 1/5] END ...C=100, gamma=0.001, kernel=poly;; score=0.862 total time=
[CV 2/5] END ...C=100, gamma=0.001, kernel=poly;; score=0.893 total time=
[CV 3/5] END ...C=100, gamma=0.001, kernel=poly;; score=0.893 total time=
[CV 4/5] END ...C=100, gamma=0.001, kernel=poly;; score=0.893 total time=
[CV 5/5] END ...C=100, gamma=0.001, kernel=poly;; score=0.893 total time=
[CV 1/5] END ....C=100, gamma=0.001, kernel=rbf;; score=0.828 total time=
[CV 2/5] END ....C=100, gamma=0.001, kernel=rbf;; score=0.893 total time=
[CV 3/5] END ....C=100, gamma=0.001, kernel=rbf;; score=0.821 total time=
[CV 4/5] END ....C=100, gamma=0.001, kernel=rbf;; score=0.929 total time=
[CV 5/5] END ....C=100, gamma=0.001, kernel=rbf;; score=0.857 total time=
[CV 1/5] END C=100, gamma=0.001, kernel=sigmoid;; score=0.828 total time=
[CV 2/5] END C=100, gamma=0.001, kernel=sigmoid;; score=0.893 total time=
```

```
#checking best parameters for linear using grid search
param_grid = {'C': [0.1, 1, 10, 100], 'kernel': ['linear']}

svml = SVC()
grid_searchl = GridSearchCV(svm, param_grid, refit=True, verbose=3)
grid_searchl.fit(X_train, y_train)

best_paramsl = grid_searchl.best_params_
best_scorel = grid_searchl.best_score_
print(f"Best Hyperparameters: {best_paramsl}")
print(f"Best Cross-Validation Accuracy: {best_scorel:.4f}")
```

```
⇒ Fitting 5 folds for each of 4 candidates, totalling 20 fits
[CV 1/5] END .....C=0.1, kernel=linear;; score=0.828 total time=
[CV 2/5] END .....C=0.1, kernel=linear;; score=0.893 total time=
[CV 3/5] END .....C=0.1, kernel=linear;; score=0.821 total time=
[CV 4/5] END .....C=0.1, kernel=linear;; score=0.893 total time=
[CV 5/5] END .....C=0.1, kernel=linear;; score=0.857 total time=
[CV 1/5] END .....C=1, kernel=linear;; score=0.862 total time=
[CV 2/5] END .....C=1, kernel=linear;; score=0.857 total time=
[CV 3/5] END .....C=1, kernel=linear;; score=0.821 total time=
[CV 4/5] END .....C=1, kernel=linear;; score=0.893 total time=
[CV 5/5] END .....C=1, kernel=linear;; score=0.857 total time=
[CV 1/5] END .....C=10, kernel=linear;; score=0.828 total time=
[CV 2/5] END .....C=10, kernel=linear;; score=0.857 total time=
[CV 3/5] END .....C=10, kernel=linear;; score=0.821 total time=
[CV 4/5] END .....C=10, kernel=linear;; score=0.929 total time=
[CV 5/5] END .....C=10, kernel=linear;; score=0.893 total time=
[CV 1/5] END .....C=100, kernel=linear;; score=0.897 total time=
[CV 2/5] END .....C=100, kernel=linear;; score=0.857 total time=
[CV 3/5] END .....C=100, kernel=linear;; score=0.857 total time=
[CV 4/5] END .....C=100, kernel=linear;; score=0.929 total time=
[CV 5/5] END .....C=100, kernel=linear;; score=0.893 total time=
Best Hyperparameters: {'C': 100, 'kernel': 'linear'}
Best Cross-Validation Accuracy: 0.8865
```

```
#performing SVM with linear kernel
svm_linear = SVC(kernel='linear', C=100)
svm_linear.fit(X_train, y_train)
y_pred_linear = svm_linear.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```

```
accuracy_linear = accuracy_score(y_test, y_pred_linear)
precision_linear = precision_score(y_test, y_pred_linear)
recall_linear = recall_score(y_test, y_pred_linear)
f1_linear = f1_score(y_test, y_pred_linear)
```

```
print("Accuracy:", accuracy_linear)
print("Precision:", precision_linear)
print("Recall:", recall_linear)
print("F1-score:", f1_linear)
```

```
➡ Accuracy: 0.9444444444444444
  Precision: 1.0
  Recall: 0.6666666666666666
  F1-score: 0.8
```

```
#checking best parameters for linear using grid search
param_grid = {'C': [0.1, 1, 10, 100], 'gamma': ['scale', 'auto', 1, 0.1, 0.01, 0.001]}
```

```
svmp = SVC()
grid_searchp = GridSearchCV(svmp, param_grid, refit=True, verbose=3)
grid_searchp.fit(X_train, y_train)
```

```
best_paramsp = grid_searchp.best_params_
best_scorep = grid_searchp.best_score_
print(f"Best Hyperparameters: {best_paramsp}")
print(f"Best Cross-Validation Accuracy: {best_scorep:.4f}")
```

```
➡ Fitting 5 folds for each of 96 candidates, totalling 480 fits
[CV 1/5] END C=0.1, degree=1, gamma=scale, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=1, gamma=scale, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=1, gamma=scale, kernel=poly;; score=0.893 total
[CV 4/5] END C=0.1, degree=1, gamma=scale, kernel=poly;; score=0.893 total
[CV 5/5] END C=0.1, degree=1, gamma=scale, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=1, gamma=auto, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=1, gamma=auto, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=1, gamma=auto, kernel=poly;; score=0.893 total
[CV 4/5] END C=0.1, degree=1, gamma=auto, kernel=poly;; score=0.893 total
```

```

[CV 5/5] END C=0.1, degree=1, gamma=auto, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=1, gamma=1, kernel=poly;; score=0.828 total tir
[CV 2/5] END C=0.1, degree=1, gamma=1, kernel=poly;; score=0.893 total tir
[CV 3/5] END C=0.1, degree=1, gamma=1, kernel=poly;; score=0.821 total tir
[CV 4/5] END C=0.1, degree=1, gamma=1, kernel=poly;; score=0.893 total tir
[CV 5/5] END C=0.1, degree=1, gamma=1, kernel=poly;; score=0.857 total tir
[CV 1/5] END C=0.1, degree=1, gamma=0.1, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=1, gamma=0.1, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=1, gamma=0.1, kernel=poly;; score=0.893 total
[CV 4/5] END C=0.1, degree=1, gamma=0.1, kernel=poly;; score=0.893 total
[CV 5/5] END C=0.1, degree=1, gamma=0.1, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=1, gamma=0.01, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=1, gamma=0.01, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=1, gamma=0.01, kernel=poly;; score=0.893 total
[CV 4/5] END C=0.1, degree=1, gamma=0.01, kernel=poly;; score=0.893 total
[CV 5/5] END C=0.1, degree=1, gamma=0.01, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=1, gamma=0.001, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=1, gamma=0.001, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=1, gamma=0.001, kernel=poly;; score=0.893 total
[CV 4/5] END C=0.1, degree=1, gamma=0.001, kernel=poly;; score=0.893 total
[CV 5/5] END C=0.1, degree=1, gamma=0.001, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=2, gamma=scale, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=2, gamma=scale, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=2, gamma=scale, kernel=poly;; score=0.857 total
[CV 4/5] END C=0.1, degree=2, gamma=scale, kernel=poly;; score=0.893 total
[CV 5/5] END C=0.1, degree=2, gamma=scale, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=2, gamma=auto, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=2, gamma=auto, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=2, gamma=auto, kernel=poly;; score=0.893 total
[CV 4/5] END C=0.1, degree=2, gamma=auto, kernel=poly;; score=0.893 total
[CV 5/5] END C=0.1, degree=2, gamma=auto, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=2, gamma=1, kernel=poly;; score=0.862 total tir
[CV 2/5] END C=0.1, degree=2, gamma=1, kernel=poly;; score=0.893 total tir
[CV 3/5] END C=0.1, degree=2, gamma=1, kernel=poly;; score=0.893 total tir
[CV 4/5] END C=0.1, degree=2, gamma=1, kernel=poly;; score=0.893 total tir
[CV 5/5] END C=0.1, degree=2, gamma=1, kernel=poly;; score=0.929 total tir
[CV 1/5] END C=0.1, degree=2, gamma=0.1, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=2, gamma=0.1, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=2, gamma=0.1, kernel=poly;; score=0.857 total
[CV 4/5] END C=0.1, degree=2, gamma=0.1, kernel=poly;; score=0.893 total
[CV 5/5] END C=0.1, degree=2, gamma=0.1, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=2, gamma=0.01, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=2, gamma=0.01, kernel=poly;; score=0.893 total
[CV 3/5] END C=0.1, degree=2, gamma=0.01, kernel=poly;; score=0.893 total
[CV 4/5] END C=0.1, degree=2, gamma=0.01, kernel=poly;; score=0.893 total
[CV 5/5] END C=0.1, degree=2, gamma=0.01, kernel=poly;; score=0.893 total
[CV 1/5] END C=0.1, degree=2, gamma=0.001, kernel=poly;; score=0.862 total
[CV 2/5] END C=0.1, degree=2, gamma=0.001, kernel=poly;; score=0.893 total

```

```
#performing SVM polynomial with best hyperparameters
svm_polynomial = SVC(kernel='poly', C=0.1, degree=2, gamma=1)
svm_polynomial.fit(X_train, y_train)
y_pred_polynomial = svm_polynomial.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```

```
accuracy_poly = accuracy_score(y_test, y_pred_polynomial)
precision_poly = precision_score(y_test, y_pred_polynomial)
recall_poly = recall_score(y_test, y_pred_polynomial)
f1_poly = f1_score(y_test, y_pred_polynomial)
```

```
print("Accuracy:", accuracy_poly)
print("Precision:", precision_poly)
print("Recall:", recall_poly)
print("F1-score:", f1_poly)
```

```
➞ Accuracy: 0.9444444444444444
Precision: 1.0
Recall: 0.6666666666666666
F1-score: 0.8
```

```
#checking best parameters for rbf using grid search
param_grid = {'C': [0.1, 1, 10, 100], 'gamma': ['scale', 'auto', 1, 0.1, 0.01, 0.001]}
```

```
svmr = SVC()
grid_searchr = GridSearchCV(svmr, param_grid, refit=True, verbose=3)
grid_searchr.fit(X_train, y_train)
```

```
best_paramsr = grid_searchr.best_params_
best_scorer = grid_searchr.best_score_
print(f"Best Hyperparameters: {best_paramsr}")
print(f"Best Cross-Validation Accuracy: {best_scorer:.4f}")
```

```
➞ Fitting 5 folds for each of 24 candidates, totalling 120 fits
[CV 1/5] END ....C=0.1, gamma=scale, kernel=rbf;; score=0.862 total time=
[CV 2/5] END ....C=0.1, gamma=scale, kernel=rbf;; score=0.893 total time=
[CV 3/5] END ....C=0.1, gamma=scale, kernel=rbf;; score=0.893 total time=
[CV 4/5] END ....C=0.1, gamma=scale, kernel=rbf;; score=0.893 total time=
[CV 5/5] END ....C=0.1, gamma=scale, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=0.1, gamma=auto, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=0.1, gamma=auto, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=0.1, gamma=auto, kernel=rbf;; score=0.893 total time=
```

```

[CV 4/5] END .....C=0.1, gamma=auto, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=0.1, gamma=auto, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 1/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.862 total time=
[CV 2/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.893 total time=
[CV 3/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.893 total time=
[CV 4/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.893 total time=
[CV 5/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=1, gamma=scale, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=1, gamma=scale, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=1, gamma=scale, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=1, gamma=scale, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=1, gamma=scale, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=1, gamma=auto, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=1, gamma=auto, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=1, gamma=auto, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=1, gamma=auto, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=1, gamma=auto, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=1, gamma=1, kernel=rbf;; score=0.897 total time=
[CV 2/5] END .....C=1, gamma=1, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=1, gamma=1, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=1, gamma=1, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=1, gamma=1, kernel=rbf;; score=0.929 total time=
[CV 1/5] END .....C=1, gamma=0.1, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=1, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=1, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=1, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=1, gamma=0.1, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=1, gamma=0.01, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=1, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 3/5] END .....C=1, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 4/5] END .....C=1, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 5/5] END .....C=1, gamma=0.01, kernel=rbf;; score=0.893 total time=
[CV 1/5] END .....C=1, gamma=0.001, kernel=rbf;; score=0.862 total time=
[CV 2/5] END .....C=1, gamma=0.001, kernel=rbf;; score=0.893 total time=

```

```
#performing SVM with RBF kernel based on best parameters
svm_rbf = SVC(kernel='rbf', C=10, gamma=0.1)
svm_rbf.fit(X_train, y_train)
y_pred_rbf = svm_polynomial.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```

```
accuracy_rbf = accuracy_score(y_test, y_pred_rbf)
precision_rbf = precision_score(y_test, y_pred_rbf)
recall_rbf = recall_score(y_test, y_pred_rbf)
f1_rbf = f1_score(y_test, y_pred_rbf)
```

```
print("Accuracy:", accuracy_rbf)
print("Precision:", precision_rbf)
print("Recall:", recall_rbf)
print("F1-score:", f1_rbf)
```

```
⇒ Accuracy: 0.9444444444444444
Precision: 1.0
Recall: 0.6666666666666666
F1-score: 0.8
```

```
#checking best parameters for sigmoid kernel using grid search
param_grid = {'C': [0.1, 1, 10, 100], 'gamma': ['scale', 'auto', 1, 0.1, 0.01, 0.001]}
```

```
svmr = SVC()
grid_searchr = GridSearchCV(svmr, param_grid, refit=True, verbose=3)
grid_searchr.fit(X_train, y_train)
```

```
best_paramsr = grid_searchr.best_params_
best_scorer = grid_searchr.best_score_
print(f"Best Hyperparameters: {best_paramsr}")
print(f"Best Cross-Validation Accuracy: {best_scorer:.4f}")
```

```
⇒ [CV 3/5] END .C=10, gamma=scale, kernel=sigmoid;; score=0.857 total time=
[CV 4/5] END .C=10, gamma=scale, kernel=sigmoid;; score=0.750 total time=
[CV 5/5] END .C=10, gamma=scale, kernel=sigmoid;; score=0.857 total time=
[CV 1/5] END ..C=10, gamma=auto, kernel=sigmoid;; score=0.828 total time=
[CV 2/5] END ..C=10, gamma=auto, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END ..C=10, gamma=auto, kernel=sigmoid;; score=0.821 total time=
[CV 4/5] END ..C=10, gamma=auto, kernel=sigmoid;; score=0.893 total time=
[CV 5/5] END ..C=10, gamma=auto, kernel=sigmoid;; score=0.857 total time=
[CV 1/5] END .....C=10, gamma=1, kernel=sigmoid;; score=0.862 total time=
```

```

[CV 2/5] END .....C=10, gamma=1, kernel=sigmoid;; score=0.857 total time=
[CV 3/5] END .....C=10, gamma=1, kernel=sigmoid;; score=0.857 total time=
[CV 4/5] END .....C=10, gamma=1, kernel=sigmoid;; score=0.857 total time=
[CV 5/5] END .....C=10, gamma=1, kernel=sigmoid;; score=0.929 total time=
[CV 1/5] END ...C=10, gamma=0.1, kernel=sigmoid;; score=0.862 total time=
[CV 2/5] END ...C=10, gamma=0.1, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END ...C=10, gamma=0.1, kernel=sigmoid;; score=0.857 total time=
[CV 4/5] END ...C=10, gamma=0.1, kernel=sigmoid;; score=0.750 total time=
[CV 5/5] END ...C=10, gamma=0.1, kernel=sigmoid;; score=0.857 total time=
[CV 1/5] END ..C=10, gamma=0.01, kernel=sigmoid;; score=0.793 total time=
[CV 2/5] END ..C=10, gamma=0.01, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END ..C=10, gamma=0.01, kernel=sigmoid;; score=0.821 total time=
[CV 4/5] END ..C=10, gamma=0.01, kernel=sigmoid;; score=0.893 total time=
[CV 5/5] END ..C=10, gamma=0.01, kernel=sigmoid;; score=0.857 total time=
[CV 1/5] END .C=10, gamma=0.001, kernel=sigmoid;; score=0.862 total time=
[CV 2/5] END .C=10, gamma=0.001, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END .C=10, gamma=0.001, kernel=sigmoid;; score=0.893 total time=
[CV 4/5] END .C=10, gamma=0.001, kernel=sigmoid;; score=0.893 total time=
[CV 5/5] END .C=10, gamma=0.001, kernel=sigmoid;; score=0.893 total time=
[CV 1/5] END C=100, gamma=scale, kernel=sigmoid;; score=0.862 total time=
[CV 2/5] END C=100, gamma=scale, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END C=100, gamma=scale, kernel=sigmoid;; score=0.857 total time=
[CV 4/5] END C=100, gamma=scale, kernel=sigmoid;; score=0.821 total time=
[CV 5/5] END C=100, gamma=scale, kernel=sigmoid;; score=0.821 total time=
[CV 1/5] END .C=100, gamma=auto, kernel=sigmoid;; score=0.828 total time=
[CV 2/5] END .C=100, gamma=auto, kernel=sigmoid;; score=0.929 total time=
[CV 3/5] END .C=100, gamma=auto, kernel=sigmoid;; score=0.786 total time=
[CV 4/5] END .C=100, gamma=auto, kernel=sigmoid;; score=0.857 total time=
[CV 5/5] END .C=100, gamma=auto, kernel=sigmoid;; score=0.857 total time=
[CV 1/5] END ....C=100, gamma=1, kernel=sigmoid;; score=0.897 total time=
[CV 2/5] END ....C=100, gamma=1, kernel=sigmoid;; score=0.821 total time=
[CV 3/5] END ....C=100, gamma=1, kernel=sigmoid;; score=0.893 total time=
[CV 4/5] END ....C=100, gamma=1, kernel=sigmoid;; score=0.750 total time=
[CV 5/5] END ....C=100, gamma=1, kernel=sigmoid;; score=0.929 total time=
[CV 1/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.862 total time=
[CV 2/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.893 total time=
[CV 4/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.821 total time=
[CV 5/5] END ..C=100, gamma=0.1, kernel=sigmoid;; score=0.857 total time=
[CV 1/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.828 total time=
[CV 2/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.821 total time=
[CV 4/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.893 total time=
[CV 5/5] END .C=100, gamma=0.01, kernel=sigmoid;; score=0.893 total time=
[CV 1/5] END C=100, gamma=0.001, kernel=sigmoid;; score=0.828 total time=
[CV 2/5] END C=100, gamma=0.001, kernel=sigmoid;; score=0.893 total time=
[CV 3/5] END C=100, gamma=0.001, kernel=sigmoid;; score=0.821 total time=
[CV 4/5] END C=100, gamma=0.001, kernel=sigmoid;; score=0.893 total time=
[CV 5/5] END C=100, gamma=0.001, kernel=sigmoid;; score=0.857 total time=

```



```
#performing svm with sigmoid kernel based on best parameters and evaluating with
svm_sigmoid = SVC(kernel='sigmoid', C=1, gamma=1)
svm_sigmoid.fit(X_train, y_train)
y_pred_sig = svm_sigmoid.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
```

```
accuracy_sig = accuracy_score(y_test, y_pred_sig)
precision_sig = precision_score(y_test, y_pred_sig)
recall_sig = recall_score(y_test, y_pred_sig)
f1_sig = f1_score(y_test, y_pred_sig)
```

```
print("Accuracy:", accuracy_sig)
print("Precision:", precision_sig)
print("Recall:", recall_sig)
print("F1-score:", f1_sig)
```

```
↪ Accuracy: 0.8611111111111112
Precision: 1.0
Recall: 0.16666666666666666
F1-score: 0.2857142857142857
```

```
#K-fold cross validation with SVM , linear kernel
from sklearn.model_selection import cross_val_score
from sklearn.svm import SVC
```

```
svm_linear = SVC(kernel='linear', C=100)
scores = cross_val_score(svm_linear, df_scaled, y, cv=5)
print("Cross-validation scores:", scores)
print("Mean accuracy:", scores.mean())
```

```
↪ Cross-validation scores: [0.86111111 0.88888889 0.88571429 0.91428571 0.8
Mean accuracy: 0.8699999999999999
```

Naïve Bayes Variant Comparison

Table 1: Performance Comparison of Naïve Bayes Variants

Metric	Gaussian NB	Multinomial NB	Bernoulli NB
Accuracy	0.8219	0.8706	0.8573
Precision	0.7233	0.8326	0.8004
Recall	0.9385	0.8547	0.8128
F1 Score	0.8170	0.8435	0.8065

KNN: Varying k Values

Table 2: KNN Performance for Different k Values

k	Accuracy	Precision	Recall	F1 Score
1	0.8958	0.8769	0.8769	0.8769
3	0.8936	0.8883	0.8564	0.8721
5	0.8958	0.9016	0.8462	0.8730
7	0.8921	0.8965	0.8423	0.8685

KNN: KDTree vs BallTree

Table 3: KNN Comparison: KDTree vs BallTree

Metric	KDTree	BallTree
Accuracy	0.8958	0.8958
Precision	0.9016	0.9016
Recall	0.8462	0.8462
F1 Score	0.8730	0.8730
Training Time (s)	0.012	0.015

SVM Performance

Table 4: SVM Performance with Different Kernels and Parameters

Kernel	Hyperparameters	Accuracy	F1 Score	Training Time (s)
Linear	$C = 100$	0.9444	0.8000	0.031
Polynomial	$C = 0.1$, degree = 1, gamma = auto	0.8930	0.8930	0.042
RBF	$C = 100$, gamma = 0.1	0.9290	0.9123	0.050
Sigmoid	$C = 100$, gamma = 0.1	0.8570	0.8502	0.047

K-Fold Cross-Validation Results

Table 5: Cross-Validation Scores for Each Model

Fold	Gaussian NB Acc.	Bernoulli NB Acc.	KNN Acc.	SVM Acc.
Fold 1	0.8261	0.8586	0.8967	0.9444
Fold 2	0.8282	0.8597	0.8956	0.9423
Fold 3	0.8173	0.8575	0.8956	0.9444
Fold 4	0.8204	0.8564	0.8945	0.9423
Fold 5	0.8204	0.8575	0.8945	0.9444
Average	0.8225	0.8579	0.8954	0.9436

Observations & Conclusions

- **Best overall accuracy:** SVM with Linear kernel ($C = 100$) achieved the highest average accuracy in K-Fold validation.
- **Best Naïve Bayes variant:** MultinomialNB had the highest single-split accuracy, though BernoulliNB was also competitive.
- **KNN trend:** Accuracy remained high for $k = 1$ to 7, with best stability at $k = 5$.
- **KDTree vs BallTree:** Both yielded identical accuracy, with KDTree slightly faster in training.
- **SVM kernels:** Linear kernel outperformed others; polynomial kernel lagged behind despite parameter tuning.